

## **Ranges-v3: A powerful library for working with ranges**

### **Introduction:**

Range-v3 is a third party open source generic library for working with range based functionalities. The functions in this library are augmented with the STL(Standard Template Library) ranges in C++. Range-v3 is designed to be compatible with the STL, and it provides implementations of many of the same algorithms and data structures. However, Range-v3 also provides a number of additional features, such as views, actions, and concepts, which make it easier to write powerful and expressive range-based code.

Range-v3 is a powerful and expressive library that can be used to write better C++ code. It is particularly well-suited for writing range-based code, such as code for machine learning, data science, and numerical computing.

I will be writing this blog based on the use cases of machine learning field such as:

- Feature engineering: Range-v3 can be used to create new features for our machine learning models. For example, you can use views to calculate statistics for your features, combine multiple features into a single feature, and create features that represent interactions between features.
- Data preprocessing: Range-v3 can be used to efficiently preprocess our data for machine learning. For example, you can use views to filter out outliers, normalize your features, and convert your data to a different format.
- Model training: Range-v3 can be used to implement custom training algorithms for our machine learning models. For example, you can use views to batch your data, shuffle your data, and evaluate our model's performance on a held-out validation set.
- Model prediction: Range-v3 can be used to make predictions with our machine learning models. For example, you can use views to batch your input data and make predictions for each batch of data.

### **Installation:**

- How to install Range-v3: Installing the library in MacOS was straightforward and WindowsOS has a bit more steps involved then MacOS.
  - Mac User: To install range-v3 in macOS, simply open the terminal using Homebrew package manager type “ brew install range-v3” and it will install the package with the most recent update.
  - Windows user: There are different ways to install the package. The following steps are for building from source.
    - Clone the Range-v3 repository:
      - **git clone <https://github.com/ericniebler/range-v3>.git**
    - Navigate to the Range-v3 directory and run the following command to generate Visual Studio project files:
      - **cmake -S . -B build -G "Visual Studio 17 2022 Win64"**
    - Open the build\Range-v3.sln solution file in Visual Studio.
    - Build the Range-v3 library by right-clicking on the Range-v3 project and selecting Build.
    - Once the Range-v3 library is built, it will be installed to the build\install\include directory.
  - Using range-v3 in our code;
    - To use Range-v3 in your code, you need to include the following header file:
      - **#include <range/v3/all.hpp>**
 This will include all of the headers.

## Compilation:

Range-v3 is a header-only library that provides a comprehensive set of algorithms and utilities for working with ranges in C++. To compile your code with range-v3, simply include the header files in your project and compile it as usual.

## Code Examples:

Let's explore how to prepare and work with data in C++ using the Range-v3 library. I've put together a series of practical examples that show how Range-v3 makes data processing easier. I'll cover basics like picking specific data, doing calculations, and organizing information. Each example comes with straightforward code and explanations to help you understand how it all works. Whether you're new to data processing or an experienced programmer, these examples will make your data tasks simpler and more efficient.

**Pipe (|):** One of the benefits of using range-v3 is that it can often make your code more concise and readable. For example, the following code shows how to double the values in a vector using the ranges STL and range-v3 libraries: `std::vector<int> foo = {1, 2, 3, 4, 5};`

```
// ranges STL
std::vector<int> bar1;
std::transform(std::begin(foo), std::end(foo), std::back_inserter(bar1), [](auto const v: const int) -> int {
    return 2 * v;
});

// range-v3
auto bar = invoke_result_t<...,> = foo | ranges::views::transform( [](auto const v: const int) -> int {
    return 2 * v;
});
```

In the above example, the precise difference between the ranges STL and range-v3 code is that the range-v3 code uses a **pipe operator (|)** to chain together the range operations. This allows you to write more concise and expressive code, without having to worry about explicit iterators. The ranges STL code uses explicit iterators to iterate over the range and apply the transformation function. This can make the code more verbose and difficult to read, especially for complex range operations.

Let's breakdown the above code for the pipe("|"):

#### Start with a Range:

- In the provided code, `foo` represents the initial range or collection of data that you want to process. It could be a container, such as a vector, containing elements.

#### Apply Operations:

- Using the pipe operator (|), you apply a series of operations to the `foo` range.
- In this specific case, the operation applied is `ranges::views::transform`, which transforms each element in the range `foo` by doubling its value.

#### Passing Results:

- The result of the `ranges::views::transform` operation is a new range (in this case, `bar`) that contains the transformed elements.
- The pipe operator allows you to pass the result of the transformation as input to subsequent operations, creating a processing pipeline.

#### Final Result:

- The final result, `bar`, is a range that holds the transformed elements. Each element in `bar` is twice the value of the corresponding element in `foo`.

The key benefit of the pipe operator is that it simplifies the code and makes it more readable.

**ranges::for\_each** is a versatile function in Range-v3 that simplifies data processing. It allows us to iterate through elements in a range while performing custom actions on each. In our example, we use it to filter the values with particular limits and print the

data elements, highlighting its power to streamline complex data tasks.

```
// using foreach and filter and print all elements in the vector.
ranges::for_each(
    rng: data | ranges::views::filter( pred: [](const auto& card: const pair<...> &) -> bool { return card.first < 10.0 && card.second < 90.0; } ),
    fun: [](const auto& card: const pair<...> &) -> void {
        std::cout << "X: " << card.first << ", Y: " << card.second << std::endl;
    });
```

The **ranges::actions::sort** function takes a 2D vector and arranges it in a specific order. It does this by looking at the first part of each pair and using that information to sort everything. This is really useful for getting your data in the right order so you can work with it more easily.

```
void SortData(std::vector<std::pair<double, double>>& data)
{
    // Sort the data based on the first value of each pair
    data |= ranges::action::sort( pred: [](const std::pair<double, double>& a, const std::pair<double, double>& b) -> bool {
        return a.first < b.first;
    });
}
```

This code snippet finds the largest element within a dataset by considering the first value in each pair. It uses **ranges::max\_element** along with a custom comparison function to determine the maximum element based on the first values of the pairs.

```
/// finds max element based on first value
auto maxElement(borrowed_iterator_t<...>) = ranges::max_element( &: data, pred: [](const std::pair<double, double>& a, const std::pair<double, double>& b) -> bool {
    return a.first < b.first;
});
```

This code snippet calculates the sum of the first values in a dataset using Range-v3. It applies the **ranges::accumulate** function to the dataset, transforming each element to extract the first value. The accumulate function then sums these first values, starting from an initial value of 0.0, providing a simple and efficient way to compute the sum of the dataset's first values.

```
/// finds max element based on first value
auto maxElement(borrowed_iterator_t<...>) = ranges::max_element( &: data, pred: [](const std::pair<double, double>& a, const std::pair<double, double>& b) -> bool {
    return a.first < b.first;
});
```

This code snippet demonstrates filtering data using Range-v3. It applies the **ranges::views::filter** function to the dataset, selecting only the elements where the first value is greater than 3.0. This simplifies the process of isolating specific data points that meet a particular criterion, improving data precision and analysis.

```
// Filtering: Filter data where the first value is greater than 3.0
auto filteredData: invoke_result_t<...> = data | ranges::views::filter( pred: [](const std::pair<double, double>& pair) -> bool {
    return pair.first > 3.0;
});
```

This code snippet showcases data mapping with Range-v3. It utilizes the **ranges::views::transform** function to square both the first and second values of each pair within the dataset. This elegant approach simplifies data transformation, allowing for the efficient and simultaneous processing of multiple data points.

```
// Mapping: Square the first and second values
auto squaredData : invoke_result_t<..., > = data | ranges::views::transform( fun: [](const std::pair<double, double>& pair) -> pair<double, double> {
    return std::make_pair( t1: pair.first * pair.first, t2: pair.second * pair.second);
});
```

This code example introduces data grouping using Range-v3. It employs the **ranges::views::group\_by** function to group data based on the first value of each pair. The provided comparison function checks if two elements have the same first value, effectively organizing data into groups where the first values match. This simplifies the process of categorizing and analyzing data with common attributes.

```
// Grouping: Group data by the first value
auto groupedData : invoke_result_t<..., > = data | ranges::views::group_by( fun: [](const std::pair<double, double>& a, const std::pair<double, double>& b) -> bool {
    return a.first == b.first;
});
```

This code snippet demonstrates searching for a specific element within a dataset using Range-v3. The **ranges::find** function is employed to locate the first element where the second value matches the pair (5.0, 6.0). This convenient feature simplifies the process of pinpointing and retrieving specific data elements within a collection

```
// Searching: Find the first element where the second value is equal to 6.0
auto foundElement : borrowed_iterator_t<...> = ranges::find( &: data, val: std::make_pair( t1: 5.0, t2: 6.0));
```

This code snippet illustrates data joining in Range-v3. It utilizes the **ranges::views::concat** function to combine two identical copies of the dataset into a single range. This straightforward approach simplifies the process of merging data, creating a unified collection for further analysis or processing.

```
// Joining: Concatenate two copies of the data into a single range
auto concatenatedData : concat_view<...> = ranges::views::concat( &: data, &: data);
```

## Recommended use:

**When to Use Range-v3:** Range-v3 shines when you want to make data-related tasks in C++ simpler and more readable. It's a great choice for transforming, analyzing, and filtering data. If you're looking to create clean and expressive code for tasks like data grouping or summation, Range-v3 is your friend. When you need to chain multiple data processing steps together in a clear way, the pipe operator in Range-v3 makes it a breeze. Plus, if you have special data processing needs, Range-v3 lets you create your own custom tools to get the job done efficiently.

**When Not to Use Range-v3:** However, Range-v3 might not be the best fit for every scenario. If you're building a performance-critical application where every bit of speed matters, you might want to consider lower-level C++ constructs to avoid any potential overhead. If you're working with older C++ standards, Range-v3, which relies on C++17 or later features, won't be an option. Small-scale projects with straightforward data processing needs may not benefit significantly from Range-v3's capabilities, and introducing it could add unnecessary complexity. Similarly, for real-time systems with strict timing constraints, avoiding Range-v3 might be a wise choice to ensure predictable performance. Lastly, if you're in an environment with limited community support or a team unfamiliar with Range-v3, it may be challenging to adopt effectively.

## Tips for new users:

**Start Small:** Begin with simple examples and gradually work your way up to more complex data processing tasks. Range-v3 can be used for a wide range of applications, so it's helpful to build a strong foundation.

**Read the Documentation:** Explore the official Range-v3 documentation, which is filled with clear examples and guidelines. It's an excellent resource to understand the library's features and capabilities.

**Learn Lambda Functions:** Understanding lambda functions in C++ is crucial, as you'll often use them to customize data processing. A basic grasp of lambda expressions allows you to personalize operations.

**Use the Pipe Operator:** The pipe operator (`|`) is a key feature in Range-v3. It simplifies your code by creating data processing pipelines, improving readability and organization.

**Optimize for Performance:** Keep performance in mind when working with Range-v3, especially for large datasets. Ensure your code remains efficient by optimizing where necessary.

**Read and Share Code:** Learning from others' code examples and sharing your own can be a valuable learning tool. It allows you to see different approaches to data processing tasks with Range-v3.

## Conclusion:

In the world of C++ data processing, the Range-v3 library offers a versatile toolkit of functions and classes designed to streamline the manipulation and transformation of data. These functions bring a touch of functional programming to C++, simplifying tasks such as filtering, mapping, and aggregation. What's particularly powerful is the ability to customize these operations using lambda functions, allowing you to tailor data processing to your specific needs. Range-v3 introduces classes like 'views,' 'actions,' and 'algorithms,' which encapsulate complex data processing logic and provide a structured approach to handling data. The pipe operator ('|') serves as the linchpin of data processing, facilitating the creation of elegant processing pipelines. Furthermore, the library's extensibility allows you to craft custom range adaptors and algorithms, perfect for domain-specific operations. Range-v3's lazy evaluation ensures efficient data processing, while error handling remains an essential consideration. Overall, the functions and classes in Range-v3 empower developers to write expressive, efficient code, making data manipulation a breeze.

I have provided my github repo link that I have used for this code snippet, it has more details if you want to explore range-v3 functions and also enclosed Range-V3 manual.

## Resources:

My Git repo for the code details: [https://github.com/Gaya1858/Blog\\_Range\\_v3](https://github.com/Gaya1858/Blog_Range_v3)

Range-V3 manual: <https://ericniebler.github.io/range-v3/>

**I may try to write below for the final blog:**

**Functions and classes:**

- Overview of the main functions and classes in Range-v3
- Your experience trying to use them
- Difficulty of using them

**Integration:**

- How to integrate Range-v3 into your existing code
- Any challenges you faced
- Difficulty of integration